

Homework 7 Written Part

Q1: Feed Forward and Backpropagation

1a) In the neural network structure, $f_1(\cdot)$ is a sigmoid function and $f_2(\cdot)$ is a softmax function. The sigmoid function is defined as below:

$$f_1(z) = \frac{1}{1 + e^{-z}}$$

The sigmoid function's derivative can be defined as below:

Through power rule and chain rule:

$$f_1 = (1 + e^{-z})^{-1} \rightarrow f_1' = -(1 + e^{-z})^{-2} \cdot -e^{-z} = \frac{e^{-z}}{(1 + e^{-z})^2}$$

We can split the denominator and manipulate the numerator as below:

$$f_1' = \frac{1}{(1 + e^{-z})} \cdot \frac{e^{-z} + 1 - 1}{(1 + e^{-z})} = \frac{1}{(1 + e^{-z})} \left(1 - \frac{1}{(1 + e^{-z})} \right)$$

Replacing $\frac{1}{1 + e^{-z}}$ with f_1 , which is the sigmoid function definition:

$$\boxed{f_1'(z) = f_1(z)(1 - f_1(z))}$$

Next, the softmax function is defined as below:

$$f_2(z_i) = \frac{e^{z_i}}{\sum_j^K e^{z_j}}$$

In the above function, the z_i can be broken down into K sub-components (z_j). We want to differentiate each and every one of these sub-components because the total derivative make up the entire derivative in respect to z_i , given us the expression below:

Through quotient rule:

$$\frac{\partial f_2}{\partial z_j} = \frac{\sum_j^K e^{z_j} \frac{\partial}{\partial z_j} e^{z_i} - e^{z_i} \frac{\partial}{\partial z_j} \sum_j^K e^{z_j}}{\left(\sum_j^K e^{z_j} \right)^2}$$

There are two underlying scenarios we want to observe: when $i = j$, and when $i \neq j$, the derivatives can be defined depending on the scenario:

$$\frac{\partial}{\partial z_j} e^{z_i} = \begin{cases} e^{z_j} & i = j \\ 0 & i \neq j \end{cases}$$

$$\frac{\partial}{\partial z_j} \sum_j^K e^{z_j} = \frac{\partial}{\partial z_j} (e^{z_j} + \dots + e^{z_K}) = e^{z_j}$$

Substituting these definition into $\frac{\partial f_2}{\partial z_j}$ starting with scenario where $i = j$:

$$\frac{\partial f_2}{\partial z_j} = \frac{e^{z_j} \sum_j^K e^{z_j} - e^{z_j} e^{z_i}}{\left(\sum_j^K e^{z_j} \right)^2} = \frac{e^{z_j} \left(\sum_j^K e^{z_j} - e^{z_i} \right)}{\left(\sum_j^K e^{z_j} \right)^2}$$

Similarly to the sigmoid function, we want to split the denominator and manipulate the numerator a little bit:

$$= \frac{e^{z_j}}{\sum_j^K e^{z_j}} \left(\frac{\sum_j^K e^{z_j}}{\sum_j^K e^{z_j}} - \frac{e^{z_i}}{\sum_j^K e^{z_j}} \right) = \frac{e^{z_j}}{\sum_j^K e^{z_j}} \left(1 - \frac{e^{z_i}}{\sum_j^K e^{z_j}} \right)$$

Replace with softmax function definition:

$$\frac{\partial f_2}{\partial z_j} = \boxed{f_2(z_j)(1 - f_2(z_i)), i = j}$$

When $i \neq j$:

$$\frac{\partial f_2}{\partial z_j} = \frac{e^{z_j} \cdot 0 - e^{z_j} e^{z_i}}{\left(\sum_j^K e^{z_j}\right)^2} = -\left(\frac{e^{z_j}}{\sum_j^K e^{z_j}}\right) \left(\frac{e^{z_i}}{\sum_j^K e^{z_j}}\right) = \boxed{-f_2(z_j)f_2(z_i), i \neq j}$$

1b) The error cost for output is defined as the cross-entropy function: $L = -t^T \ln(x_2)$. To define the partial derivative of L w.r.t. x_2 :

$$\frac{\partial L}{\partial x_2} = -t^T \frac{\partial}{\partial x_2} \ln(x_2) = \boxed{-t^T \frac{1}{x_2}}$$

1c) Because a_1, x_1 and δ_1 are not at the output layer, they have the following definition:

$$a_i = W_i \cdot x_{i-1} + b_i \quad x_i = f_i(a_i) \quad \delta_i = \delta_{i+1} W_{i+1} f'_i(a_i)$$

Using the provided information such as:

$$W_1 = \begin{bmatrix} 1 & 2 & -3 & 0 & 1 & -3 \\ 3 & 1 & 2 & 1 & 0 & 2 \\ 2 & 2 & 2 & 2 & 2 & 1 \\ 1 & 0 & 2 & 1 & -2 & 2 \end{bmatrix} \quad W_2 = \begin{bmatrix} 1 & 2 & -2 & 1 \\ 1 & -1 & 1 & 2 \\ 3 & 1 & -1 & 1 \end{bmatrix} \quad x_0 = [1, 1, 0, 0, 1, 1]^T \quad t = [0, 1, 0]^T \quad b_i = 1, \forall i$$

$$a_1 = W_1 x_0 + b_1 = \begin{bmatrix} 1 & 2 & -3 & 0 & 1 & -3 \\ 3 & 1 & 2 & 1 & 0 & 2 \\ 2 & 2 & 2 & 2 & 2 & 1 \\ 1 & 0 & 2 & 1 & -2 & 2 \end{bmatrix}_{4 \times 6} \begin{bmatrix} 1 \\ 1 \\ 0 \\ 0 \\ 1 \\ 1 \end{bmatrix}_{6 \times 1} + \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \end{bmatrix}_{4 \times 1} = \boxed{\begin{bmatrix} 2 \\ 7 \\ 8 \\ 2 \end{bmatrix}_{4 \times 1}}$$

$$x_1 = f_1(a_1) = \begin{bmatrix} \frac{1}{1+e^{-2}} \\ \frac{1}{1+e^{-7}} \\ \frac{1}{1+e^{-8}} \\ \frac{1}{1+e^{-2}} \end{bmatrix} = \boxed{\begin{bmatrix} 0.8808 \\ 0.9991 \\ 0.9997 \\ 0.8808 \end{bmatrix}}$$

Since $\delta_1 = \delta_2 W_2 f'_1(a_1)$ and it contains components (δ_2) at the output layer, we need a different set of definition to describe those components. Note that variables need to cross over to Exercise 1d in order to solve for δ_1 in Exercise 1c, Exercise 1d starts here to prevent repeated work and I will put answer for δ_1 at the end of Exercise 1d.

1d)

$$\delta_2 = \frac{\partial L}{\partial x_2} \frac{\partial x_2}{\partial a_2} = \frac{\partial L}{\partial x_2} \circ f'_2(a_2)$$

$$a_2 = W_2 \cdot x_1 + b_2 = \begin{bmatrix} 1 & 2 & -2 & 1 \\ 1 & -1 & 1 & 2 \\ 3 & 1 & -1 & 1 \end{bmatrix}_{3 \times 4} \begin{bmatrix} 0.8808 \\ 0.9991 \\ 0.9997 \\ 0.8808 \end{bmatrix}_{4 \times 1} + \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix}_{3 \times 1} = \boxed{\begin{bmatrix} 2.7604 \\ 3.6430 \\ 4.5226 \end{bmatrix}_{3 \times 1}}$$

$$x_2 = f_2(a_2) = f_2(W_2 x_1) = \frac{e^{W_2 x_1}}{\sum_{j=1}^{K=3} e^{W_{2,j} x_1}} = \frac{e^{W_2 x_1}}{e^{W_{2,1} x_1} + e^{W_{2,2} x_1} + e^{W_{2,3} x_1}}$$

where $W_{2,j}$ refers to the j th row in W_2 . Continuing to evaluate:

$$x_2 = \frac{e^{a_2}}{e^{2.8763} + e^{3.6430} + e^{4.5226}} = \boxed{\begin{bmatrix} 0.1082 \\ 0.2615 \\ 0.6303 \end{bmatrix}}$$

To match the matrix multiplication, the order of multiplication is switched, such that:

$$\delta_2 = \frac{\partial x_2}{\partial a_2} \frac{\partial L}{\partial x_2} = -f'_2(a_2) \circ t^T \frac{1}{x_2} = -f'_2 \left(\begin{bmatrix} 2.8763 \\ 3.6430 \\ 4.5226 \end{bmatrix} \right)_{3 \times 3} \left(\begin{bmatrix} 0/0.1082 \\ 1/0.2615 \\ 0/0.6303 \end{bmatrix} \right)_{3 \times 1}$$

f'_2 is a 3x3 matrix from due to the nature of the derivative of the softmax function, the 3x3 matrix can be represented as below, where $a_{2,j}$ represents to the j th row in a_2 :

$$f'_2(a_2) = \begin{bmatrix} f_2(a_{2,1})(1 - f_2(a_{2,1})) & -f_2(a_{2,1})f_2(a_{2,2}) & -f_2(a_{2,1})f_2(a_{2,3}) \\ -f_2(a_{2,2})f_2(a_{2,1}) & f_2(a_{2,2})(1 - f_2(a_{2,2})) & -f_2(a_{2,2})f_2(a_{2,3}) \\ -f_2(a_{2,3})f_2(a_{2,1}) & -f_2(a_{2,3})f_2(a_{2,2}) & f_2(a_{2,3})(1 - f_2(a_{2,3})) \end{bmatrix}$$

$$f'_2(a_2) = \begin{bmatrix} x_{2,1}(1 - x_{2,1}) & -x_{2,1}x_{2,2} & -x_{2,1}x_{2,3} \\ -x_{2,2}x_{2,1} & x_{2,2}(1 - x_{2,2}) & -x_{2,2}x_{2,3} \\ -x_{2,3}x_{2,1} & -x_{2,3}x_{2,2} & x_{2,3}(1 - x_{2,3}) \end{bmatrix}$$

$$= \begin{bmatrix} 0.1082(1 - 0.1082) & -(0.1082)(0.2615) & -(0.1082)(0.6303) \\ -(0.2615)(0.1082) & 0.2615(1 - 0.2615) & -(0.2615)(0.6303) \\ -(0.6303)(0.1082) & -(0.6303)(0.2615) & 0.6303(1 - 0.6303) \end{bmatrix} = \begin{bmatrix} 0.0965 & -0.0283 & -0.0682 \\ -0.0283 & 0.1931 & -0.1648 \\ -0.0682 & -0.1648 & 0.2330 \end{bmatrix}$$

Tying the terms altogether:

$$\delta_2 = - \begin{bmatrix} 0.0965 & -0.0283 & -0.0682 \\ -0.0283 & 0.1931 & -0.1648 \\ -0.0682 & -0.1648 & 0.2330 \end{bmatrix} \begin{bmatrix} 0 \\ 3.8236 \\ 0 \end{bmatrix} = \begin{bmatrix} 0.1082 \\ -0.7385 \\ 0.6303 \end{bmatrix}$$

Going back to part 1c to δ_1 , where $\delta_1 = \delta_2 W_2 f'_1(a_1)$ We are still missing the value $f'_1(a_1)$, therefore solving as follows:

$$f'_1(a_1) = f_1(a_1)(1 - f_1(a_1)) = [x_{1,1}(1 - x_{1,1}) \ x_{1,2}(1 - x_{1,2}) \ x_{1,3}(1 - x_{1,3}) \ x_{1,4}(1 - x_{1,4})]^T$$

$$= \begin{bmatrix} 0.1050 \\ 0.0009 \\ 0.0003 \\ 0.1050 \end{bmatrix}$$

We need to modify the multiplication sequence a little bit in order to utilize the Hadamard Product, δ_1 becomes $W^T \delta_2 \circ f'_1(a_1)$:

$$\delta_1 = \begin{bmatrix} 1 & 1 & 3 \\ 2 & -1 & 1 \\ -2 & 1 & -1 \\ 1 & 2 & 1 \end{bmatrix}_{4 \times 3} \begin{bmatrix} 0.1082 \\ -0.7385 \\ 0.6303 \end{bmatrix}_{3 \times 1} \circ \begin{bmatrix} 0.1050 \\ 0.0009 \\ 0.0003 \\ 0.1050 \end{bmatrix}_{4 \times 1} = \begin{bmatrix} 0.1323 \\ 0.0014 \\ -0.0005 \\ -0.0775 \end{bmatrix}$$

1e) Based on the result from 1d, the output of x_2 from the input x_0 equals to $[0.1082, 0.2615, 0.6303]^T$, therefore, the index with maximum probability will be chosen: $[0, 0, 1]^T$, which corresponds to class 3.

1f) The loss function is defined by the cross entropy function, in which the loss can be found as below:

$$L = -t^T \ln(x_2) = -[0 \ 1 \ 0] \begin{bmatrix} \ln(0.1082) \\ \ln(0.2615) \\ \ln(0.6303) \end{bmatrix} = 1.3413$$

2a) The gradient descent of W_2 is defined by:

$$W_{2,new} \leftarrow W_{2,old} - \alpha \frac{\partial L}{\partial w_2}$$

This can be further defined as:

$$W_{2,new} \leftarrow W_{2,old} - \alpha \delta_2 x_1$$

For a learning rate of 0.5, the updated W_2 is calculated as below, with x_1 transposed to match the matrix dimension for subtraction:

$$W_{2,new} = \begin{bmatrix} 1 & 2 & -2 & 1 \\ 1 & -1 & 1 & 2 \\ 3 & 1 & -1 & 1 \end{bmatrix} - 0.5 \begin{bmatrix} 0.1082 \\ -0.7385 \\ 0.6303 \end{bmatrix} \begin{bmatrix} 0.8808 & 0.9991 & 0.9997 & 0.8808 \end{bmatrix} =$$

$$\begin{bmatrix} 0.9524 & 1.9460 & -2.0541 & 0.9524 \\ 1.3252 & -0.6311 & 1.3691 & 2.3252 \\ 2.7224 & 0.6852 & -1.3150 & 0.7224 \end{bmatrix}$$

2b) Define gradient descent update for b_2 as:

$$b_{2,new} \leftarrow b_{2,old} - \alpha \delta_2$$

$$b_{2,new} = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} - 0.5 \begin{bmatrix} 0.1082 \\ -0.7385 \\ 0.6303 \end{bmatrix} = \begin{bmatrix} 0.9459 \\ 1.3692 \\ 0.6849 \end{bmatrix}$$

2c)

$$W_{1,new} \leftarrow W_{1,old} - \alpha \delta_1 x_0$$

$$W_{1,new} = \begin{bmatrix} 1 & 2 & -3 & 0 & 1 & -3 \\ 3 & 1 & 2 & 1 & 0 & 2 \\ 2 & 2 & 2 & 2 & 2 & 1 \\ 1 & 0 & 2 & 1 & -2 & 2 \end{bmatrix} - 0.5 \begin{bmatrix} 0.1323 \\ 0.0014 \\ -0.0005 \\ -0.0775 \end{bmatrix} \begin{bmatrix} 1 & 1 & 0 & 0 & 1 & 1 \end{bmatrix} =$$

$$\begin{bmatrix} 0.9338 & 1.9338 & -3.0000 & 0 & 0.9338 & -3.0662 \\ 2.9993 & 0.9993 & 2.0000 & 1.0000 & -0.0007 & 1.9993 \\ 2.0003 & 2.0003 & 2.0000 & 2.0000 & 2.0003 & 1.0003 \\ 1.0388 & 0.0388 & 2.0000 & 1.0000 & -1.9612 & 2.0388 \end{bmatrix}$$

2d)

$$b_{1,new} \leftarrow b_{1,old} - \alpha \delta_1$$

$$b_{1,new} = \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \end{bmatrix} - 0.5 \begin{bmatrix} 0.1323 \\ 0.0014 \\ -0.0005 \\ -0.0775 \end{bmatrix} = \begin{bmatrix} 0.9338 \\ 0.9993 \\ 1.0003 \\ 1.0388 \end{bmatrix}$$

2e) The updated x_2 using forward propagation of updated components is equal to $\begin{bmatrix} 0.0517, 0.8566, 0.0918 \end{bmatrix}$, which corresponds to a classification of $[0, 1, 0]$, same as the ground truth.

Q2

4/2/2021

q2

```
In [11]: import numpy as np
import matplotlib.pyplot as plt
from sklearn import tree, neighbors, linear_model, svm, neural_network, metrics
```

Problem 2a

Generating Dataset

```
In [2]: X = np.genfromtxt('winequality-red.csv', delimiter=';')
```

```
In [3]: train_X = X[1:1001, 0:-1]
train_Y = np.array([X[1:1001, -1]>=6]).astype(int).squeeze()
test_X = X[1001::, 0:-1]
test_Y = np.array([X[1001::, -1]>=6]).astype(int).squeeze()
```

Decision Tree

```
In [4]: metrics = ["gini", "entropy"]
max_depth = [3, 5, 8]
for i in metrics:
    for j in max_depth:
        clf = tree.DecisionTreeClassifier(criterion=i, max_depth=int(j)).fit(train_X, train_Y)
        train_preds = clf.predict(train_X)
        test_preds = clf.predict(test_X)
        train_acc = clf.score(train_X, train_Y)*100
        test_acc = clf.score(test_X, test_Y)*100
        print("Criterion: {}, Max Depth: {}, Training Accuracy: {}, Test Accuracy: {}.\n".format(i, j, train_acc, test_acc))
```

Criterion: gini, Max Depth: 3, Training Accuracy: 73.8, Test Accuracy: 71.61936560934892.

Criterion: gini, Max Depth: 5, Training Accuracy: 79.4, Test Accuracy: 72.95492487479132.

Criterion: gini, Max Depth: 8, Training Accuracy: 89.1, Test Accuracy: 67.9465776293823.

Criterion: entropy, Max Depth: 3, Training Accuracy: 73.4, Test Accuracy: 71.78631051752922.

Criterion: entropy, Max Depth: 5, Training Accuracy: 77.7, Test Accuracy: 73.78964941569282.

Criterion: entropy, Max Depth: 8, Training Accuracy: 86.9, Test Accuracy: 69.44908180300501.

K-Nearest Neighbors

```
In [5]: n = [8, 12, 25]
for i in n:
    clf = neighbors.KNeighborsClassifier(n_neighbors=i).fit(train_X, train_Y)
```

file:///C:/Users/ivanw/Documents/College & Grad School/Carnegie Mellon/24-787 Machine Learning/HW7/q2.html

1/6

Q2

4/2/2021

q2

```
train_preds = clf.predict(train_X)
test_preds = clf.predict(test_X)
train_acc = clf.score(train_X, train_Y)*100
test_acc = clf.score(test_X, test_Y)*100
print("N-neighbors: {}, Training Accuracy: {}, Test Accuracy: {}".format(i, train_acc, test_acc))
```

N-neighbors: 8, Training Accuracy: 75.5, Test Accuracy: 58.931552587646074.

N-neighbors: 12, Training Accuracy: 74.1, Test Accuracy: 59.59933222036727.

N-neighbors: 25, Training Accuracy: 71.8, Test Accuracy: 62.604340567612695.

Logistic Regression

In [6]:

```
train_X_wbias = np.hstack((train_X, np.ones_like(train_Y)[: , np.newaxis]))
test_X_wbias = np.hstack((test_X, np.ones_like(test_Y)[: , np.newaxis]))
reg_method = ["l1", "l2", "elasticnet"]
for i in reg_method:
    if i == "l1":
        clf = linear_model.LogisticRegression(penalty=i, solver="saga", max_iter=100000)
    elif i == "elasticnet":
        clf = linear_model.LogisticRegression(penalty=i, solver="saga", l1_ratio=1, max_iter=100000)
    else:
        clf = linear_model.LogisticRegression(penalty=i, max_iter=100000).fit(train_X_wbias, train_Y)
    train_preds = clf.predict(train_X_wbias)
    test_preds = clf.predict(test_X_wbias)
    train_acc = clf.score(train_X_wbias, train_Y)*100
    test_acc = clf.score(test_X_wbias, test_Y)*100
    print("Penalty: {}, Training Accuracy: {}, Test Accuracy: {}".format(i, train_acc, test_acc))
```

Penalty: l1, Training Accuracy: 72.8, Test Accuracy: 75.45909849749583.

Penalty: l2, Training Accuracy: 73.7, Test Accuracy: 76.62771285475793.

Penalty: elasticnet, Training Accuracy: 72.8, Test Accuracy: 75.45909849749583.

Support Vector Machine

In [7]:

```
kernels = ["linear", "rbf"]
Cs = [0.01, 10, 1000]
for i in kernels:
    for j in Cs:
        clf = svm.SVC(C=j, kernel=i).fit(train_X, train_Y)
        train_preds = clf.predict(train_X)
        test_preds = clf.predict(test_X)
        train_acc = clf.score(train_X, train_Y)*100
        test_acc = clf.score(test_X, test_Y)*100
        print("Kernel: {}, C: {}, Training Accuracy: {}, Test Accuracy: {}".format(i, j, train_acc, test_acc))
    clf = svm.SVC(kernel="poly").fit(train_X, train_Y)
    print("Kernel: poly, Training Accuracy: {}, Test Accuracy: {}".format(train_acc, test_acc))
```

Kernel: linear, C: 0.01, Training Accuracy: 72.0, Test Accuracy: 72.45409015025042.

Kernel: linear, C: 10, Training Accuracy: 73.3, Test Accuracy: 76.62771285475793.

Kernel: linear, C: 1000, Training Accuracy: 73.2, Test Accuracy: 75.79298831385643.

Q2

4/2/2021

q2

Kernel: rbf, C: 0.01, Training Accuracy: 62.3, Test Accuracy: 60.60100166944908.

Kernel: rbf, C: 10, Training Accuracy: 71.5, Test Accuracy: 71.28547579298832.

Kernel: rbf, C: 1000, Training Accuracy: 77.5, Test Accuracy: 74.95826377295492.

Kernel: poly, Training Accuracy: 77.5, Test Accuracy: 74.95826377295492.

Neural Network

```
In [8]: act_fun = ["logistic", "tanh", "relu"]
for i in act_fun:
    clf = neural_network.MLPClassifier(hidden_layer_sizes=(10, 20, 10), activation=i, m
    train_preds = clf.predict(train_X_wbias)
    test_preds = clf.predict(test_X_wbias)
    train_acc = clf.score(train_X_wbias, train_Y)*100
    test_acc = clf.score(test_X_wbias, test_Y)*100
    print("Set 1: Activation Function: {}, Training Accuracy: {}, Test Accuracy: {}.\n"
for i in act_fun:
    clf = neural_network.MLPClassifier(hidden_layer_sizes=(20, 50, 10), activation=i, m
    train_preds = clf.predict(train_X_wbias)
    test_preds = clf.predict(test_X_wbias)
    train_acc = clf.score(train_X_wbias, train_Y)*100
    test_acc = clf.score(test_X_wbias, test_Y)*100
    print("Set 2: Activation Function: {}, Training Accuracy: {}, Test Accuracy: {}.\n"
for i in act_fun:
    clf = neural_network.MLPClassifier(hidden_layer_sizes=(30, 70, 10), activation=i, m
    train_preds = clf.predict(train_X_wbias)
    test_preds = clf.predict(test_X_wbias)
    train_acc = clf.score(train_X_wbias, train_Y)*100
    test_acc = clf.score(test_X_wbias, test_Y)*100
    print("Set 3: Activation Function: {}, Training Accuracy: {}, Test Accuracy: {}.\n"
```

Set 1: Activation Function: logistic, Training Accuracy: 75.0, Test Accuracy: 75.79298831385643.

Set 1: Activation Function: tanh, Training Accuracy: 75.8, Test Accuracy: 73.12186978297161.

Set 1: Activation Function: relu, Training Accuracy: 74.4, Test Accuracy: 73.62270450751252.

Set 2: Activation Function: logistic, Training Accuracy: 74.7, Test Accuracy: 76.12687813021702.

Set 2: Activation Function: tanh, Training Accuracy: 76.6, Test Accuracy: 71.28547579298832.

Set 2: Activation Function: relu, Training Accuracy: 76.4, Test Accuracy: 72.78797996661102.

Set 3: Activation Function: logistic, Training Accuracy: 74.7, Test Accuracy: 75.79298831385643.

Set 3: Activation Function: tanh, Training Accuracy: 76.4, Test Accuracy: 72.45409015025042.

Set 3: Activation Function: relu, Training Accuracy: 75.4, Test Accuracy: 77.46243739565

Q2

4/2/2021

q2

944.

Problem 2b

A confusion matrix is a matrix that takes a prediction i and try to categorize it with to a ground truth j , if i and j does not match, it is a false positive or negative, otherwise, it is a true positive and negative.

A precision takes components from confusion matrix (true positives and false positives) and put it in a ratio such that it represents a performance/efficacy value for the prediction algorithm.

A recall matrix components acts as a precision, but on the true positive and false negative from the confusion matrix. It is a measurement of the ability of the prediction algorithm to output positive values.

A F1 matrix is the weighted average between the recall and precision.

Best parameters from each category

Decision Tree: Criterion = entropy, Max Depth = 5, Training Accuracy: 77.7, Test Accuracy: 73.45575959933221.

```
In [12]: clf = tree.DecisionTreeClassifier(criterion="entropy", max_depth=5).fit(train_X, train_Y)
test_preds = clf.predict(test_X)
confusion = metrics.confusion_matrix(test_Y, test_preds)
precision = metrics.precision_score(test_Y, test_preds)
recall = metrics.recall_score(test_Y, test_preds)
f1 = metrics.f1_score(test_Y, test_preds)
print("Confusion Matrix =\n", confusion)
print("Precision = ", precision)
print("Recall = ", recall)
print("F1 = ", f1)
```

```
Confusion Matrix =
[[145  94]
 [ 65 295]]
Precision = 0.7583547557840618
Recall = 0.8194444444444444
F1 = 0.787716955941255
```

KNN: N-neighbors = 25, Training Accuracy: 71.8, Test Accuracy: 62.604340567612695.

```
In [13]: clf = neighbors.KNeighborsClassifier(n_neighbors=25).fit(train_X, train_Y)
test_preds = clf.predict(test_X)
confusion = metrics.confusion_matrix(test_Y, test_preds)
precision = metrics.precision_score(test_Y, test_preds)
recall = metrics.recall_score(test_Y, test_preds)
f1 = metrics.f1_score(test_Y, test_preds)
print("Confusion Matrix =\n", confusion)
print("Precision = ", precision)
print("Recall = ", recall)
print("F1 = ", f1)
```


Q2

4/2/2021

q2

```
Confusion Matrix =  
[[133 106]  
 [118 242]]  
Precision = 0.6954022988505747  
Recall = 0.6722222222222223  
F1 = 0.6836158192090396
```

Logistic Regression: Penalty = l2, Training Accuracy: 73.7, Test Accuracy: 76.62771285475793.

```
In [14]: clf = linear_model.LogisticRegression(max_iter=10000).fit(train_X_wbias, train_Y)  
test_preds = clf.predict(test_X_wbias)  
confusion = metrics.confusion_matrix(test_Y, test_preds)  
precision = metrics.precision_score(test_Y, test_preds)  
recall = metrics.recall_score(test_Y, test_preds)  
f1 = metrics.f1_score(test_Y, test_preds)  
print("Confusion Matrix =\n", confusion)  
print("Precision = ", precision)  
print("Recall = ", recall)  
print("F1 = ", f1)
```

```
Confusion Matrix =  
[[156 83]  
 [ 57 303]]  
Precision = 0.7849740932642487  
Recall = 0.8416666666666667  
F1 = 0.8123324396782843
```

Support Vector Machine: Kernel = linear, C = 10, Training Accuracy: 73.3, Test Accuracy: 76.62771285475793.

```
In [15]: clf = svm.SVC(C=10, kernel="linear").fit(train_X, train_Y)  
test_preds = clf.predict(test_X)  
confusion = metrics.confusion_matrix(test_Y, test_preds)  
precision = metrics.precision_score(test_Y, test_preds)  
recall = metrics.recall_score(test_Y, test_preds)  
f1 = metrics.f1_score(test_Y, test_preds)  
print("Confusion Matrix =\n", confusion)  
print("Precision = ", precision)  
print("Recall = ", recall)  
print("F1 = ", f1)
```

```
Confusion Matrix =  
[[170 69]  
 [ 71 289]]  
Precision = 0.8072625698324022  
Recall = 0.8027777777777778  
F1 = 0.8050139275766017
```

Neural Network: Hidden Layer Sizes = 20, 50, 10, Activation Function = Logistic

```
In [16]: clf = neural_network.MLPClassifier(hidden_layer_sizes=(20, 50, 10), activation=i, max_i  
train_preds = clf.predict(train_X_wbias)  
test_preds = clf.predict(test_X_wbias)  
confusion = metrics.confusion_matrix(test_Y, test_preds)  
precision = metrics.precision_score(test_Y, test_preds)  
recall = metrics.recall_score(test_Y, test_preds)
```

Q2

4/2/2021

q2

```
f1 = metrics.f1_score(test_Y, test_preds)
print("Confusion Matrix =\n", confusion)
print("Precision = ", precision)
print("Recall = ", recall)
print("F1 = ", f1)
```

```
Confusion Matrix =
[[159  80]
 [ 77 283]]
Precision = 0.7796143250688705
Recall = 0.7861111111111111
F1 = 0.7828492392807745
```

Case study

Case A: I would sell the Logistic Regression model in this case because I want to maximize the number of positive value (good wine), therefore, I should look for the model with the highest Recall Score.

Case B: I would sell the Logistic Regression model again in this case. The customer does not want to miss too many good wines, this implies that she wants a good recall score on the good wines. However, she also does not desire when classified good wines to actually taste bad (false positives), which implies a precision scoring. Together, I conclude that the best F1 score will satisfy her the most, which the Logistic Regression model is able to offer.